


State-Driven Retrieval and Learned Re-Ranking for Conversational Music Recommendation

Nidhin Pattaniyil · Semih Yagli · Tanwir Zaman
Independent Researchers, Hoboken, New Jersey, USA

11
retrieval branches

146
ranker features

29/40
where that finished



What the system reads off a conversation

User: “...subdued, female singer, a sense of place in the lyrics, stark almost a cappella delivery... not it either.”

EXTRACTED STATE

anchor artist	Neko Case (must use)
rejected	“Calling Cards”, “Man”
facets	mood = subdued · theme = sense of place sonic = stark, almost a cappella
request	hidden_target


The full schema: request type · intent mode · anchors and pinned tracks · per-track feedback · rejected tracks and artists · hard filters · mood, genre and sonic facets · process constraints · routing flags · lyrical theme · release-year window.

What the ranker leans on

FEATURE FAMILY	SHARE OF GAIN
Bi-encoder cosine	59.0%
Per-branch rank and score	11.5%
Session and artist affinity	10.7%
Popularity · state · tags	9.2%
Other similarity cosines	8.9%
Constraints and rejections	0.7%

LightGBM LambdaMART · 146 features per (turn, track) pair.

THE BI-ENCODER BEHIND THAT TOP FEATURE



How it scored

OFFICIAL SPLIT	NDCG@20	JUDGE	COMPOSITE
Blind-A · development phase	0.4380	4.70	0.5799
Blind-B · final, 29th of 40	0.2537	3.30	0.3811

What is not ideal

We tuned on what we measured.

The development set was both our training and our evaluation data, and the response prompt was picked against the Blind-A leaderboard: 4.70 there, 3.30 on Blind-B.

We extracted more than we used.

Most fields in the state never reached retrieval or ranking at all.

Nothing from outside the catalog.

No outside source for what a track or artist is known for, and no tempo, key or soundtrack. Past one CF vector, no listening behaviour.

The reply was never checked.

Asked for a track we do not have, it described another by the same artist.

Asked for, served

the Kamelot track from Silverthorn	→	Kamelot – Fallen Star, from Haven CONSTRAINT NEVER APPLIED
exactly 126.70 bpm , in G minor	→	no tempo or key field exists to match on NOTHING IN THE CATALOG
“Watercolors” by Pat Metheny	→	Metheny – Alfie, described as if it were RESPONSE LLM


The target often ignores the request

User: “...similar chill vibe, but **from a different artist?**” (just played: Bonobo)
Target: **Bonobo – “Jets”** · annotated MOVES toward goal
Next turn: “...hoping for something **from a different artist this time...**”
Target: **Bonobo – “Pieces”** · annotated MOVES toward goal

Ground truth is whatever played next in an LLM-generated session, so we had every turn re-scored.

HOW THE LLM RELABELLING WORKED

Both judges score the same two axes: different artist asked for? does the track fit?



61%

of training turns where the judges would not settle on the target being a good fit

17%

keep the artist the listener had just rejected



Paper, code, trained models and data
github.com/npatta01/music-crs-2026