

RECSYS CHALLENGE 2026 • MUSIC-CRS

State-Driven Retrieval and Learned Re-Ranking for Conversational Music Recommendation

Nidhin Pattaniyil • Semih Yagli • Tanwir Zaman

47,071


tracks in the catalog

11

ways to search

1

learned ranker





The short version

SUMMARY

WHAT WE BUILT

An LLM reads each turn into a **typed state**; eleven retrieval branches search one index; a **LightGBM** ranker picks 20; an LLM writes the reply.

HOW IT SCORED

0.3811

Blind-B · 29th of 40
nDCG@20 fell from 0.438 on Blind-A to 0.254.

WHAT WE FOUND

31 of 43


diagnosed failures were ranking

Extraction was **3**. The state was richer than what the ranker used.

The system understood the conversation more often than it **acted on it**.

≡ One state, eleven branches, one learned ranker

APPROACH



Organisers supplied the audio (CLAP), cover-art (SigLIP-2), listening-data (BPR) and small text (Qwen3-0.6B) embeddings; we added Qwen3-8B text embeddings and our own bi-encoder.

Request understanding

APPROACH

WE HAD JUST PLAYED



Ryan Adams & The Cardinals — “Let Us Down Easy”

LISTENER

...more tracks with that kind of emotional depth and acoustic focus, maybe branching out to other singer-songwriter or alt-country artists?

LLM



REQUEST TYPE	new artist
AVOID	Ryan Adams & The Cardinals
MOOD	introspective, emotional depth
SOUND	acoustic-driven
GENRE	singer-songwriter, alt-country

The LLM reads **every turn so far**: “other artists” after a Ryan Adams track becomes an artist to avoid, and the mood carries over.

Retrieval — eleven branches

APPROACH



×1

BM25 + tag resolver

keywords and resolved tags
from the state



×3

Dense text

Qwen3 embeddings of lyrics,
metadata, attributes



×1

CLAP text → audio

the sonic description, matched
to audio



×3

Anchor centroids

near liked tracks: cover art,
audio, listening



×1

User-CF centroid

this listener's own BPR vector



×2

Catalog lookups

the artist's discography; era
popularity

Each branch returns its **top 500**; the ranker sees their union.



What the ranker leans on

APPROACH



Share of the ranker's split gain, by feature family. One signal — the bi-encoder cosine — carries **59%**; constraints and rejections carry **0.7%**.

✓ How it scored

RESULTS

OFFICIAL CODABENCH SPLIT	NDCG@20	CAT.	LEX.	JUDGE	COMPOSITE
Blind-A · development phase	0.4380	0.0313	0.8028	4.70	0.5799
Blind-B · final	0.2537	0.0315	0.7862	3.30	0.3811

NDCG@20

-42%

0.4380 → 0.2537

JUDGE (1-5)

-1.40

4.70 → 3.30

DIVERSITY TERMS

flat

catalog and reply-wording
diversity barely moved

Blind-B also withholds the conversation-goal annotations — but the deployed ranker never used them at inference, so that is not the explanation.



Extracted, but the ranker did not act on it

WENT WRONG



*"I gotta break away from **Deltron** and **Del** completely"*

Dead Prez — Hip Hop

state held both artists as avoid

NOT ACTED ON



*"branching out to **other** singer-songwriter artists"*

Ryan Adams — When The Stars Go Blue

state held Ryan Adams as avoid

NOT ACTED ON



*"still focusing on **Juanes**... haven't found the depth I'm after"*

Juanes — Me Enamorado De Ti

state held Juanes as rejected

NOT ACTED ON

We wanted the ranker to use what the state recorded; it served that artist anyway. We think the labels taught it to: they often repeat the just-played artist, so rejections carried only **0.7%** of its decisions.

Where the failures were — all 80 Blind-B turns

ANALYSIS

DIAGNOSED GAP	TURNS	MEANING
Ranking	31	a better, non-violating candidate sat below our top-20
Constraint not applied	8	an extracted filter never reached the pool
State extraction	3	the extractor missed something the user said
Entity resolution	1	a name never resolved to a catalog ID

TOP-1 FLAGGED

17 / 80

the track we led with violated something the state had recorded

HOW THIS WAS COUNTED

43 diagnoses across 32 distinct turns. Our audit of each turn's system trace. Blind-B has no public answers.

Extraction is **3 of 43** diagnoses. The failure is overwhelmingly downstream of understanding the conversation.



Asked for, served

WENT WRONG

1

*“the Kamelot track from **Silverthorn**”*

Kamelot — Fallen Star

from *Haven*. We extracted “Silverthorn” but never filtered on it.

CONSTRAINT

2

“exactly 126.70 bpm, in G minor”

no tempo or key in the catalog

and we added no data from outside it

NO DATA

3

“‘Watercolors’ by Pat Metheny”

Pat Metheny — Alfie

reply: “Got it — here’s ‘Alfie’...” — “Watercolors” is not in the catalog

RESPONSE



What the training labels look like

WENT WRONG



"I keep asking for Arcade Fire — that's the only song I want"

label: Arctic Monkeys — This House is a Circus

DOES NOT



"similar chill vibe, but from a different artist"

label: Bonobo — Jets

MOVES



"what other artists fit this darkwave style?"

label: Tropic of Cancer — The Dull Age

MOVES



"what other orchestral tracks would you recommend?"

label: Nick Phoenix — Kogan

MOVES

Three of four are annotated **MOVES TOWARD GOAL** — the dataset records them as having helped. Train on this and the model learns to anchor. **Ours does.**

Why: the answer is the simulated recommender's pick from a 16–32-track pool, made without seeing the listener's goal.



Everything is released

THANK YOU



Paper, poster, slides, code, models, data

npatta01.github.io/music-crs-2026